

Spring Data REST - Reference Documentation

Jon Brisbin, Oliver Gierke, Greg Turnquist

Version 2.6.10.RELEASE, 2018-01-24

Table of Contents

Preface	2
Project metadata	3
1. Dependencies	4
1.1. Dependency management with Spring Boot	5
1.2. Spring Framework	5
Reference Documentation	6
2. Introduction	7
3. Getting started	8
3.1. Introduction	8
3.2. Adding Spring Data REST to a Spring Boot project	8
3.3. Adding Spring Data REST to a Gradle project	8
3.4. Adding Spring Data REST to a Maven project	9
3.5. Configuring Spring Data REST	9
3.6. Basic settings for Spring Data REST	9
3.6.1. Which repositories get exposed by defaults?	9
3.6.2. Changing the base URI	10
3.6.3. Changing other Spring Data REST properties	11
3.7. Starting the application	11
4. Repository resources	13
4.1. Fundamentals	13
4.1.1. Default status codes	13
4.1.2. Resource discoverability	13
4.2. The collection resource	14
4.2.1. Supported HTTP Methods	14
4.3. The item resource	15
4.3.1. Supported HTTP methods	15
4.4. The association resource	17
4.4.1. Supported HTTP methods	17
4.5. The search resource	18
4.5.1. Supported HTTP methods	18
4.6. The query method resource	18
4.6.1. Supported HTTP methods	18
5. Paging and Sorting	20
5.1. Paging	20
5.1.1. Previous and Next Links	20
5.2. Sorting	22
6. Domain Object Representations	24
6.1. Object Mapping	24

6.1.1. Adding custom (de)serializers to Jackson's ObjectMapper	24
7. Projections and Excerpts	26
7.1. Projections	26
7.1.1. Finding existing projections	29
7.1.2. Bringing in hidden data	30
7.2. Excerpts	31
7.3. Excerpting commonly accessed data	32
8. Conditional Operations with Headers	34
8.1. ETag, If-Match, and If-None-Match headers	34
8.2. If-Modified-Since header	35
8.3. Architecting a more efficient frontend	36
9. Validation	37
9.1. Assigning Validators manually	37
10. Events	38
10.1. Writing an ApplicationListener	38
10.2. Writing an annotated handler	38
11. Integration	40
11.1. Programmatic Links	40
12. Metadata	42
12.1. Application-Level Profile Semantics (ALPS)	42
12.1.1. Hypermedia control types	45
12.1.2. ALPS with Projections	46
12.1.3. Adding custom details to your ALPS descriptions	48
12.2. JSON Schema	50
13. Security	53
13.1. @Pre and @Post security	53
13.2. @Secured security	54
13.3. Enabling method level security	54
14. Tools	56
14.1. The HAL Browser	56
15. Customizing Spring Data REST	60
15.1. Customizing item resource URIs	60
15.2. Configuring the REST URL path	61
15.2.1. Handling rels	62
15.2.2. Hiding certain repositories, query methods, or fields	64
15.2.3. Hiding repository CRUD methods	65
15.3. Adding Spring Data REST to an existing Spring MVC Application	66
15.3.1. More on required configuration	66
15.4. Overriding Spring Data REST Response Handlers	67
15.5. Customizing the JSON output	69
15.5.1. The ResourceProcessor interface	69

15.5.2. Adding Links	70
15.5.3. Customizing the representation	70
15.6. Adding custom (de)serializers to Jackson's ObjectMapper	70
15.6.1. Abstract class registration	70
15.6.2. Adding custom serializers for domain types	71
15.7. Configuring CORS	72
15.7.1. Repository interface CORS configuration	72
15.7.2. Repository REST Controller method CORS configuration	73
15.7.3. Global CORS configuration	73
Appendix	74
Appendix A: Using cURL to talk to Spring Data REST	75
Appendix B: Spring Data REST example projects	76
Multi-store example	76
Projections	76
Spring Data REST + Spring Security	76
Starbucks example	76

NOTE

Copies of this document may be made for your own use and for distribution to others, provided that you do not charge any fee for such copies and further provided that each copy contains this Copyright Notice, whether distributed in print or electronically.

Preface

Project metadata

- Version control - <https://github.com/spring-projects/spring-data-rest>
- Bugtracker - <https://jira.spring.io/browse/DATAREST>
- Project page - <http://projects.spring.io/spring-data-rest>
- Release repository - <https://repo.spring.io/libs-release>
- Milestone repository - <https://repo.spring.io/libs-milestone>
- Snapshot repository - <https://repo.spring.io/libs-snapshot>

Chapter 1. Dependencies

Due to different inception dates of individual Spring Data modules, most of them carry different major and minor version numbers. The easiest way to find compatible ones is by relying on the Spring Data Release Train BOM we ship with the compatible versions defined. In a Maven project you'd declare this dependency in the `<dependencyManagement />` section of your POM:

Example 1. Using the Spring Data release train BOM

```
<dependencyManagement>
  <dependencies>
    <dependency>
      <groupId>org.springframework.data</groupId>
      <artifactId>spring-data-releasetrain</artifactId>
      <version>${release-train}</version>
      <scope>import</scope>
      <type>pom</type>
    </dependency>
  </dependencies>
</dependencyManagement>
```

The current release train version is **Ingalls-SR10**. The train names are ascending alphabetically and currently available ones are listed [here](#). The version name follows the following pattern: `${name}-${release}` where release can be one of the following:

- **BUILD-SNAPSHOT** - current snapshots
- **M1**, **M2** etc. - milestones
- **RC1**, **RC2** etc. - release candidates
- **RELEASE** - GA release
- **SR1**, **SR2** etc. - service releases

A working example of using the BOMs can be found in our [Spring Data examples repository](#). If that's in place declare the Spring Data modules you'd like to use without a version in the `<dependencies />` block.

Example 2. Declaring a dependency to a Spring Data module

```
<dependencies>
  <dependency>
    <groupId>org.springframework.data</groupId>
    <artifactId>spring-data-jpa</artifactId>
  </dependency>
</dependencies>
```


1.1. Dependency management with Spring Boot

Spring Boot already selects a very recent version of Spring Data modules for you. In case you want to upgrade to a newer version nonetheless, simply configure the property `spring-data-releasetrain.version` to the `train name and iteration` you'd like to use.

1.2. Spring Framework

The current version of Spring Data modules require Spring Framework in version 4.3.14.RELEASE or better. The modules might also work with an older bugfix version of that minor version. However, using the most recent version within that generation is highly recommended.

Reference Documentation

Chapter 2. Introduction

REST web services have become the number one means for application integration on the web. In its core, REST defines that a system consists of resources that clients interact with. These resources are implemented in a hypermedia driven way. Spring MVC offers a solid foundation to build these kinds of services. But implementing even the simplest tenet of REST web services for a multi-domain object system can be quite tedious and result in a lot of boilerplate code.

Spring Data REST builds on top of Spring Data repositories and automatically exports those as REST resources. It leverages hypermedia to allow clients to find functionality exposed by the repositories and integrates these resources into related hypermedia based functionality automatically.

Chapter 3. Getting started

3.1. Introduction

Spring Data REST is itself a Spring MVC application and is designed in such a way that it should integrate with your existing Spring MVC applications with very little effort. An existing (or future) layer of services can run alongside Spring Data REST with only minor considerations.

3.2. Adding Spring Data REST to a Spring Boot project

The simplest way to get to started is if you are building a Spring Boot application. That's because Spring Data REST has both a starter as well as auto-configuration.

Spring Boot configuration with Gradle

```
dependencies {  
    ...  
    compile("org.springframework.boot:spring-boot-starter-data-rest")  
    ...  
}
```

Spring Boot configuration with Maven

```
<dependencies>  
    ...  
    <dependency>  
        <groupId>org.springframework.boot</groupId>  
        <artifactId>spring-boot-starter-data-rest</artifactId>  
    </dependency>  
    ...  
</dependencies>
```

NOTE

You don't have to supply the version number if you are using the [Spring Boot Gradle plugin](#) or the [Spring Boot Maven plugin](#).

When using Spring Boot, Spring Data REST gets configured automatically.

3.3. Adding Spring Data REST to a Gradle project

To add Spring Data REST to a Gradle-based project, add the `spring-data-rest-webmvc` artifact to your compile-time dependencies:

```
dependencies {  
    ... other project dependencies  
    compile("org.springframework.data:spring-data-rest-webmvc:2.6.10.RELEASE")  
}
```

3.4. Adding Spring Data REST to a Maven project

To add Spring Data REST to a Maven-based project, add the `spring-data-rest-webmvc` artifact to your compile-time dependencies:

```
<dependency>  
  <groupId>org.springframework.data</groupId>  
  <artifactId>spring-data-rest-webmvc</artifactId>  
  <version>2.6.10.RELEASE</version>  
</dependency>
```

3.5. Configuring Spring Data REST

To install Spring Data REST alongside your existing Spring MVC application, you need to include the appropriate MVC configuration. Spring Data REST configuration is defined in a class called `RepositoryRestMvcConfiguration` and that class can just be imported into your applications configuration.

IMPORTANT

This step is unnecessary if you are using Spring Boot's auto-configuration. Spring Boot will automatically enable Spring Data REST when you include `spring-boot-starter-data-rest` and either in your list of dependencies, and you your app is flagged with either `@SpringBootApplication` or `@EnableAutoConfiguration`.

To customize the configuration, register a `RepositoryRestConfigurer` (or extend `RepositoryRestConfigurerAdapter`) and implement or override the `configure...`-methods relevant to your use case.

Make sure you also configure Spring Data repositories for the store you use. For details on that, please consult the reference documentation for the [corresponding Spring Data module](#).

3.6. Basic settings for Spring Data REST

3.6.1. Which repositories get exposed by defaults?

Spring Data REST uses a `RepositoryDetectionStrategy` to determine if a repository will be exported as REST resource or not. The following strategies (enumeration values of `RepositoryDiscoveryStrategies`) are available:

Table 1. Repository detection strategies

Name	Description
DEFAULT	Exposes all public repository interfaces but considers <code>@(Repository)RestResource</code> 's <code>exported</code> flag.
ALL	Exposes all repositories independently of type visibility and annotations.
ANNOTATION	Only repositories annotated with <code>@(Repository)RestResource</code> are exposed, unless their <code>exported</code> flag is set to <code>false</code> .
VISIBILITY	Only public repositories annotated are exposed.

3.6.2. Changing the base URI

By default, Spring Data REST serves up REST resources at the root URI, `"/`". There are multiple ways to change the base path.

With Spring Boot 1.2+, all it takes is a single property in `application.properties`:

```
spring.data.rest.basePath=/api
```

With Spring Boot 1.1 or earlier, or if you are not using Spring Boot, simply do this:

```
@Configuration
class CustomRestMvcConfiguration {

    @Bean
    public RepositoryRestConfigurer repositoryRestConfigurer() {

        return new RepositoryRestConfigurerAdapter() {

            @Override
            public void configureRepositoryRestConfiguration(RepositoryRestConfiguration
config) {
                config.setBasePath("/api");
            }
        };
    }
}
```

Alternatively just register a custom implementation of `RepositoryRestConfigurer` as Spring bean and make sure it gets picked up by component scanning:

```

@Component
public class CustomizedRestMvcConfiguration extends RepositoryRestConfigurerAdapter {

    @Override
    public void configureRepositoryRestConfiguration(RepositoryRestConfiguration config)
    {
        config.setBasePath("/api");
    }
}

```

Both of these approaches will change the base path to `/api`.

3.6.3. Changing other Spring Data REST properties

There are many properties you can alter:

Table 2. Spring Boot configurable properties

Name	Description
basePath	root URI for Spring Data REST
defaultPageSize	change default number of items served in a single page
maxPageSize	change maximum number of items in a single page
pageParamName	change name of the query parameter for selecting pages
limitParamName	change name of the query parameter for number of items to show in a page
sortParamName	change name of the query parameter for sorting
defaultMediaType	change default media type to use when none is specified
returnBodyOnCreate	change if a body should be returned on creating a new entity
returnBodyOnUpdate	change if a body should be returned on updating an entity

3.7. Starting the application

At this point, you must also configure your key data store.

Spring Data REST officially supports:

- [Spring Data JPA](#)
- [Spring Data MongoDB](#)
- [Spring Data Neo4j](#)

- [Spring Data GemFire](#)
- [Spring Data Cassandra](#)

Here are some Getting Started guides to help you get up and running quickly:

- [Spring Data JPA](#)
- [Spring Data MongoDB](#)
- [Spring Data Neo4j](#)
- [Spring Data GemFire](#)

These linked guides introduce how to add dependencies for the related data store, configure domain objects, and define repositories.

You can run your application as either a Spring Boot app (with links shown above) or configure it as a classic Spring MVC app.

NOTE

In general Spring Data REST doesn't add functionality to a given data store. This means that by definition, it should work with any Spring Data project that supports the Repository programming model. The data stores listed above are simply the ones we have written integration tests to verify.

From this point, you can be free to [customize Spring Data REST](#) with various options.

Chapter 4. Repository resources

4.1. Fundamentals

The core functionality of Spring Data REST is to export resources for Spring Data repositories. Thus, the core artifact to look at and potentially tweak to customize the way the exporting works is the repository interface. Assume the following repository interface:

```
public interface OrderRepository extends CrudRepository<Order, Long> { }
```

For this repository, Spring Data REST exposes a collection resource at `/orders`. The path is derived from the uncapitalized, pluralized, simple class name of the domain class being managed. It also exposes an item resource for each of the items managed by the repository under the URI template `/orders/{id}`.

By default the HTTP methods to interact with these resources map to the according methods of `CrudRepository`. Read more on that in the sections on [collection resources](#) and [item resources](#).

4.1.1. Default status codes

For the resources exposed, we use a set of default status codes:

- **200 OK** - for plain `GET` requests.
- **201 Created** - for `POST` and `PUT` requests that create new resources.
- **204 No Content** - for `PUT`, `PATCH`, and `DELETE` requests if the configuration is set to not return response bodies for resource updates (`RepositoryRestConfiguration.returnBodyOnUpdate`). If the configuration value is set to include responses for `PUT`, **200 OK** will be returned for updates, **201 Created** will be returned for resource created through `PUT`.

If the configuration values (`RepositoryRestConfiguration.returnBodyOnUpdate` and `RepositoryRestConfiguration.returnBodyCreate`) are explicitly set to null, the presence of the HTTP Accept header will be used to determine the response code.

4.1.2. Resource discoverability

A core principle of [HATEOAS](#) is that resources should be discoverable through the publication of links that point to the available resources. There are a few competing de-facto standards of how to represent links in JSON. By default, Spring Data REST uses [HAL](#) to render responses. HAL defines links to be contained in a property of the returned document.

Resource discovery starts at the top level of the application. By issuing a request to the root URL under which the Spring Data REST application is deployed, the client can extract a set of links from the returned JSON object that represent the next level of resources that are available to the client.

For example, to discover what resources are available at the root of the application, issue an HTTP `GET` to the root URL:

```
curl -v http://localhost:8080/

< HTTP/1.1 200 OK
< Content-Type: application/hal+json

{ "_links" : {
  "orders" : {
    "href" : "http://localhost:8080/orders"
  },
  "profile" : {
    "href" : "http://localhost:8080/api/alps"
  }
}
```

The property of the result document is an object in itself consisting of keys representing the relation type with nested link objects as specified in HAL.

NOTE

For more details about the **profile** link, see [Application-Level Profile Semantics \(ALPS\)](#).

4.2. The collection resource

Spring Data REST exposes a collection resource named after the uncapitalized, pluralized version of the domain class the exported repository is handling. Both the name of the resource and the path can be customized using the `@RepositoryRestResource` on the repository interface.

4.2.1. Supported HTTP Methods

Collections resources support both `GET` and `POST`. All other HTTP methods will cause a `405 Method Not Allowed`.

GET

Returns all entities the repository servers through its `findAll(...)` method. If the repository is a paging repository we include the pagination links if necessary and additional page metadata.

Parameters

If the repository has pagination capabilities the resource takes the following parameters:

- `page` - the page number to access (0 indexed, defaults to 0).
- `size` - the page size requested (defaults to 20).
- `sort` - a collection of sort directives in the format `($propertyname,)+[asc|desc]?`.

Custom status codes

- **405 Method Not Allowed** - if the `findAll(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

Supported media types

- `application/hal+json`
- `application/json`

Related resources

- `search` - a [search resource](#) if the backing repository exposes query methods.

HEAD

Returns whether the collection resource is available.

POST

Creates a new entity from the given request body.

Custom status codes

- **405 Method Not Allowed** - if the `save(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

Supported media types

- `application/hal+json`
- `application/json`

4.3. The item resource

Spring Data REST exposes a resource for individual collection items as sub-resources of the collection resource.

4.3.1. Supported HTTP methods

Item resources generally support `GET`, `PUT`, `PATCH` and `DELETE` unless explicit configuration prevents that (see below for details).

GET

Returns a single entity.

Custom status codes

- **405 Method Not Allowed** - if the `findOne(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

Supported media types

- application/hal+json
- application/json

Related resources

For every association of the domain type we expose links named after the association property. This can be customized by using `@RestResource` on the property. The related resources are of type [association resource](#).

HEAD

Returns whether the item resource is available.

PUT

Replaces the state of the target resource with the supplied request body.

Custom status codes

- **405 Method Not Allowed** - if the `save(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

Supported media types

- application/hal+json
- application/json

PATCH

Similar to **PUT** but partially updating the resources state.

Custom status codes

- **405 Method Not Allowed** - if the `save(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

Supported media types

- application/hal+json
- application/json
- [application/patch+json](#)
- [application/merge-patch+json](#)

DELETE

Deletes the resource exposed.

Custom status codes

- **405 Method Not Allowed** - if the `delete(...)` methods was not exported (through `@RestResource(exported = false)`) or is not present in the repository at all.

4.4. The association resource

Spring Data REST exposes sub-resources of every item resource for each of the associations the item resource has. The name and path of the resource defaults to the name of the association property and can be customized using `@RestResource` on the association property.

4.4.1. Supported HTTP methods

GET

Returns the state of the association resource

Supported media types

- `application/hal+json`
- `application/json`

PUT

Binds the resource pointed to by the given URI(s) to the resource. This

Custom status codes

- **400 Bad Request** - if multiple URIs were given for a to-one-association.

Supported media types

- `text/uri-list` - URIs pointing to the resource to bind to the association.

POST

Only supported for collection associations. Adds a new element to the collection.

Supported media types

- `text/uri-list` - URIs pointing to the resource to add to the association.

DELETE

Unbinds the association.

Custom status codes

- **405 Method Not Allowed** - if the association is non-optional.

4.5. The search resource

The search resource returns links for all query methods exposed by a repository. The path and name of the query method resources can be modified using `@RestResource` on the method declaration.

4.5.1. Supported HTTP methods

As the search resource is a read-only resource it supports `GET` only.

GET

Returns a list of links pointing to the individual query method resources

Supported media types

- `application/hal+json`
- `application/json`

Related resources

For every query method declared in the repository we expose a [query method resource](#). If the resource supports pagination, the URI pointing to it will be a URI template containing the pagination parameters.

HEAD

Returns whether the search resource is available. A 404 return code indicates no query method resources available at all.

4.6. The query method resource

The query method resource executes the query exposed through an individual query method on the repository interface.

4.6.1. Supported HTTP methods

As the search resource is a read-only resource it supports `GET` only.

GET

Returns the result of the query execution.

Parameters

If the query method has pagination capabilities (indicated in the URI template pointing to the resource) the resource takes the following parameters:

- `page` - the page number to access (0 indexed, defaults to 0).

- **size** - the page size requested (defaults to 20).
- **sort** - a collection of sort directives in the format (**\$propertyname**,)+[**asc|desc**]?

Supported media types

- application/hal+json
- application/json

HEAD

Returns whether a query method resource is available.

Chapter 5. Paging and Sorting

This documents Spring Data REST's usage of the Spring Data Repository paging and sorting abstractions. To familiarize yourself with those features, please see the Spring Data documentation for the Repository implementation you're using.

5.1. Paging

Rather than return everything from a large result set, Spring Data REST recognizes some URL parameters that will influence the page size and starting page number.

If you extend `PagingAndSortingRepository<T, ID>` and access the list of all entities, you'll get links to the first 20 entities. To set the page size to any other number, add a `size` parameter:

```
http://localhost:8080/people/?size=5
```

This will set the page size to 5.

To use paging in your own query methods, you need to change the method signature to accept an additional `Pageable` parameter and return a `Page` rather than a `List`. For example, the following query method will be exported to `/people/search/nameStartsWith` and will support paging:

```
@RestResource(path = "nameStartsWith", rel = "nameStartsWith")
public Page findByNameStartsWith(@Param("name") String name, Pageable p);
```

The Spring Data REST exporter will recognize the returned `Page` and give you the results in the body of the response, just as it would with a non-paged response, but additional links will be added to the resource to represent the previous and next pages of data.

5.1.1. Previous and Next Links

Each paged response will return links to the previous and next pages of results based on the current page using the IANA defined link relations `prev` and `next`. If you are currently at the first page of results, however, no `prev` link will be rendered. The same is true for the last page of results: no `next` link will be rendered.

Look at the following example, where we set the page size to 5:

```
curl localhost:8080/people?size=5
```



```
{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons{&sort,page,size}", ①
      "templated" : true
    },
    "next" : {
      "href" : "http://localhost:8080/persons?page=1&size=5{&sort}", ②
      "templated" : true
    }
  },
  "_embedded" : {
    ... data ...
  },
  "page" : { ③
    "size" : 5,
    "totalElements" : 50,
    "totalPages" : 10,
    "number" : 0
  }
}
```

At the top, we see `_links`:

- ① This `self` link serves up the whole collection with some options
- ② This `next` link points to the next page, assuming the same page size.
- ③ At the bottom is extra data about the page settings, including the size of a page, total elements, total pages, and the page number you are currently viewing.

NOTE

When using tools like **curl** on the command line, if you have a "&" in your statement, wrap the whole URI with quotes.

It's also important to notice that the `self` and `next` URIs are, in fact, URI templates. They accept not only `size`, but also `page`, `sort` as optional flags.

As mentioned, at the bottom of the HAL document, is a collection of details about the page. This extra information makes it very easy for you to configure UI tools like sliders or indicators to reflect overall position the user is in viewing the data. For example, the document above shows we are looking at the first page (with page numbers indexed to 0 being the first).

What happens if we follow the `next` link?

```
$ curl "http://localhost:8080/persons?page=1&size=5"
```

```

{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons{&sort,projection,page,size}",
      "templated" : true
    },
    "next" : {
      "href" : "http://localhost:8080/persons?page=2&size=5{&sort,projection}", ①
      "templated" : true
    },
    "prev" : {
      "href" : "http://localhost:8080/persons?page=0&size=5{&sort,projection}", ②
      "templated" : true
    }
  },
  "_embedded" : {
    ... data ...
  },
  "page" : {
    "size" : 5,
    "totalElements" : 50,
    "totalPages" : 10,
    "number" : 1 ③
  }
}

```

This looks very similar, except for the following differences:

- ① The `next` link now points to yet another page, indicating it's relative perspective to the `self` link.
- ② A `prev` link now appears, giving us a path to the previous page.
- ③ The current number is now 1 (indicating the second page).

This feature makes it quite easy to map optional buttons on the screen to these hypermedia controls, hence allowing easy navigational features for the UI experience without having to hard code the URIs. In fact, the user can be empowered to pick from a list of page sizes, dynamically changing the content served, without having to rewrite the `next` and `prev` controls at the top or bottom.

5.2. Sorting

Spring Data REST recognizes sorting parameters that will use the Repository sorting support.

To have your results sorted on a particular property, add a `sort` URL parameter with the name of the property you want to sort the results on. You can control the direction of the sort by appending a `,` to the the property name plus either `asc` or `desc`. The following would use the `findByNameStartsWith` query method defined on the `PersonRepository` for all `Person` entities with names starting with the letter "K" and add sort data that orders the results on the `name` property in descending order:

```
curl -v "http://localhost:8080/people/search/nameStartsWith?name=K&sort=name,desc"
```

To sort the results by more than one property, keep adding as many `sort=PROPERTY` parameters as you need. They will be added to the `Pageable` in the order they appear in the query string. Results can be sorted by top-level and nested properties. Use property path notation to express a nested sort property. Sorting by linkable associations (i.e. resources to top-level resources) is not supported.

Chapter 6. Domain Object Representations

6.1. Object Mapping

Spring Data REST returns a representation of a domain object that corresponds to the requested **Accept** type specified in the HTTP request.

Currently, only JSON representations are supported. Other representation types can be supported in the future by adding an appropriate converter and updating the controller methods with the appropriate content-type.

Sometimes the behavior of the Spring Data REST's **ObjectMapper**, which has been specially configured to use intelligent serializers that can turn domain objects into links and back again, may not handle your domain model correctly. There are so many ways one can structure your data that you may find your own domain model isn't being translated to JSON correctly. It's also sometimes not practical in these cases to try and support a complex domain model in a generic way. Sometimes, depending on the complexity, it's not even possible to offer a generic solution.

6.1.1. Adding custom (de)serializers to Jackson's **ObjectMapper**

To accommodate the largest percentage of use cases, Spring Data REST tries very hard to render your object graph correctly. It will try and serialize unmanaged beans as normal POJOs and it will try and create links to managed beans where that's necessary. But if your domain model doesn't easily lend itself to reading or writing plain JSON, you may want to configure Jackson's **ObjectMapper** with your own custom type mappings and (de)serializers.

Abstract class registration

One key configuration point you might need to hook into is when you're using an abstract class (or an interface) in your domain model. Jackson won't know by default what implementation to create for an interface. Take the following example:

```
@Entity
public class MyEntity {
    @OneToMany
    private List<MyInterface> interfaces;
}
```

In a default configuration, Jackson has no idea what class to instantiate when POSTing new data to the exporter. This is something you'll need to tell Jackson either through an annotation, or, more cleanly, by registering a type mapping using a **Module**.

To add your own Jackson configuration to the **ObjectMapper** used by Spring Data REST, override the **configureJacksonObjectMapper** method. That method will be passed an **ObjectMapper** instance that has a special module to handle serializing and deserializing `PersistentEntity`s`. You can register your own modules as well, like in the following example.

```

@Override
protected void configureJacksonObjectMapper(ObjectMapper objectMapper) {
    objectMapper.registerModule(new SimpleModule("MyCustomModule") {
        @Override
        public void setupModule(SetupContext context) {
            context.addAbstractTypeResolver(
                new SimpleAbstractTypeResolver().addMapping(MyInterface.class,
                                                            MyInterfaceImpl.class)
            );
        }
    });
}

```

Once you have access to the `SetupContext` object in your `Module`, you can do all sorts of cool things to configure Jackson's JSON mapping. You can read more about how `Modules` work on Jackson's wiki: <http://wiki.fasterxml.com/JacksonFeatureModules>

Adding custom serializers for domain types

If you want to (de)serialize a domain type in a special way, you can register your own implementations with Jackson's `ObjectMapper` and the Spring Data REST exporter will transparently handle those domain objects correctly. To add serializers, from your `setupModule` method implementation, do something like the following:

```

@Override
public void setupModule(SetupContext context) {
    SimpleSerializers serializers = new SimpleSerializers();
    SimpleDeserializers deserializers = new SimpleDeserializers();

    serializers.addSerializer(MyEntity.class, new MyEntitySerializer());
    deserializers.addDeserializer(MyEntity.class, new MyEntityDeserializer());

    context.addSerializers(serializers);
    context.addDeserializers(deserializers);
}

```

Chapter 7. Projections and Excerpts

Spring Data REST presents a default view of the domain model you are exporting. But sometimes, you may need to alter the view of that model for various reasons. In this section, you will learn how to define projections and excerpts to serve up simplified and reduced views of resources.

7.1. Projections

Look at the following domain model:

```
@Entity
public class Person {

    @Id @GeneratedValue
    private Long id;
    private String firstName, lastName;

    @OneToOne
    private Address address;
    ...
}
```

This **Person** has several attributes:

- **id** is the primary key
- **firstName** and **lastName** are data attributes
- **address** is a link to another domain object

Now assume we create a corresponding repository as follows:

```
interface PersonRepository extends CrudRepository<Person, Long> {}
```

By default, Spring Data REST will export this domain object including all of its attributes. **firstName** and **lastName** will be exported as the plain data objects that they are. There are two options regarding the **address** attribute. One option is to also define a repository for **Address** objects like this:

```
interface AddressRepository extends CrudRepository<Address, Long> {}
```

In this situation, a **Person** resource will render the **address** attribute as a URI to its corresponding **Address** resource. If we were to look up "Frodo" in the system, we could expect to see a HAL document like this:

```
{
  "firstName" : "Frodo",
  "lastName" : "Baggins",
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons/1"
    },
    "address" : {
      "href" : "http://localhost:8080/persons/1/address"
    }
  }
}
```

There is another route. If the `Address` domain object does not have its own repository definition, Spring Data REST will inline the data fields right inside the `Person` resource.

```
{
  "firstName" : "Frodo",
  "lastName" : "Baggins",
  "address" : {
    "street": "Bag End",
    "state": "The Shire",
    "country": "Middle Earth"
  },
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons/1"
    }
  }
}
```

But what if you don't want `address` details at all? Again, by default, Spring Data REST will export all its attributes (except the `id`). You can offer the consumer of your REST service an alternative by defining one or more projections.

```
@Projection(name = "noAddresses", types = { Person.class }) ①
interface NoAddresses { ②

    String getFirstName(); ③

    String getLastName(); ④
}
```

This projection has the following details:

- ① The `@Projection` annotation flags this as a projection. The `name` attribute provides the name of the projection, which you'll see how to use shortly. The `types` attribute targets this projection to

only apply to `Person` objects.

- ② It's a Java interface making it declarative.
- ③ It exports the `firstName`.
- ④ It exports the `lastName`.

The `NoAddresses` projection only has getters for `firstName` and `lastName` meaning that it won't serve up any address information. Assuming you have a separate repository for `Address` resources, the default view from Spring Data REST is slightly different as shown below:

```
{
  "firstName" : "Frodo",
  "lastName" : "Baggins",
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons/1{?projection}", ①
      "templated" : true ②
    },
    "address" : {
      "href" : "http://localhost:8080/persons/1/address"
    }
  }
}
```

- ① There is a new option listed for this resource, `{?projection}`.
- ② The `self` URI is a URI Template.

To view apply the projection to the resource, look up <http://localhost:8080/persons/1?projection=noAddresses>.

NOTE

The value supplied to the `projection` query parameter is the same as specified in `@Projection(name = "noAddress")`. It has nothing to do with the name of the projection's interface.

It's possible to have multiple projections.

NOTE

Visit [Projections](#) to see an example project you can experiment with.

How does Spring Data REST finds projection definitions?

- Any `@Projection` interface found in the same package as your entity definitions (or one of its sub-packages) is registered.
- You can manually register via `RepositoryRestConfiguration.getProjectionConfiguration().addProjection(...)`.

In either situation, the interface with your projection MUST have the `@Projection` annotation.

7.1.1. Finding existing projections

Spring Data REST exposes [Application-Level Profile Semantics \(ALPS\)](#) documents, a micro metadata format. To view the ALPS metadata, follow the [profile](#) link exposed by the root resource. If you navigate down to the ALPS document for [Person](#) resources (which would be [/alps/persons](#)), you can find many details about [Person](#) resources. Projections will be listed along with the details about the [GET](#) REST transition, something like this:

```
{ ...
  "id" : "get-person", ①
  "name" : "person",
  "type" : "SAFE",
  "rt" : "#person-representation",
  "descriptors" : [ {
    "name" : "projection", ②
    "doc" : {
      "value" : "The projection that shall be applied when rendering the response.
Acceptable values available in nested descriptors.",
      "format" : "TEXT"
    },
    "type" : "SEMANTIC",
    "descriptors" : [ {
      "name" : "noAddresses", ③
      "type" : "SEMANTIC",
      "descriptors" : [ {
        "name" : "firstName", ④
        "type" : "SEMANTIC"
      }, {
        "name" : "lastName", ④
        "type" : "SEMANTIC"
      } ]
    } ]
  } ]
},
...
```

- ① This part of the ALPS document shows details about [GET](#) and [Person](#) resources.
- ② Further down are the [projection](#) options.
- ③ Further down you can see projection [noAddresses](#) listed.
- ④ The actual attributes served up by this projection include [firstName](#) and [lastName](#).

NOTE

Projection definitions will be picked up and made available for clients if they are:

- Flagged with the [@Projection](#) annotation and located in the same package (or sub-package) of the domain type, OR
- Manually [registered](#) via [RepositoryRestConfiguration.getProjectionConfiguration\(\).addProjection\(...\)](#).

7.1.2. Bringing in hidden data

So far, you have seen how projections can be used to reduce the information that is presented to the user. Projections can also bring in normally unseen data. For example, Spring Data REST will ignore fields or getters that are marked up with `@JsonIgnore` annotations. Look at the following domain object:

```
@Entity
public class User {

    @Id @GeneratedValue
    private Long id;
    private String name;

    @JsonIgnore private String password; ①

    private String[] roles;
    ...
}
```

① Jackson's `@JsonIgnore` is used to prevent the `password` field from getting serialized into JSON.

This `User` class can be used to store user information as well as integration with Spring Security. If you create a `UserRepository`, the `password` field would normally have been exported. Not good! In this example, we prevent that from happening by applying Jackson's `@JsonIgnore` on the `password` field.

NOTE

Jackson will also not serialize the field into JSON if `@JsonIgnore` is on the field's corresponding getter function.

However, projections introduce the ability to still serve this field. It's possible to create a projection like this:

```
@Projection(name = "passwords", types = { User.class })
interface PasswordProjection {

    String getPassword();
}
```

If such a projection is created and used, it will side step the `@JsonIgnore` directive placed on `User.password`.

IMPORTANT

This example may seem a bit contrived, but it's possible with a richer domain model and many projections, to accidentally leak such details. Since Spring Data REST cannot discern the sensitivity of such data, it is up to the developers to avoid such situations.

Projections can also generate virtual data. Imagine you had the following entity definition:

```
@Entity
public class Person {

    ...
    private String firstName;
    private String lastName;

    ...
}
```

You can create a projection that combines these two data fields together like this:

```
@Projection(name = "virtual", types = { Person.class })
public interface VirtualProjection {

    @Value("#{target.firstName} #{target.lastName}") ①
    String getFullName();

}
```

① Spring's `@Value` annotation lets you plugin a SpEL expression that takes the target object, and splices together its `firstName` and `lastName` attributes to render a read-only `fullName`.

7.2. Excerpts

An excerpt is a projection that is applied to a resource collection automatically. For an example, you can alter the `PersonRepository` as follows:

```
@RepositoryRestResource(excerptProjection = NoAddresses.class)
interface PersonRepository extends CrudRepository<Person, Long> {}
```

This directs Spring Data REST to use the `NoAddresses` projection when embedding `Person` resources into collections or related resources.

NOTE

Excerpt projections are **NOT** applied to *single resources* automatically. They have to be applied deliberately. Excerpt projections are meant to provide a default preview of collection data, but not when fetching individual resources. See [Why is an excerpt projection not applied automatically for a Spring Data REST item resource?](#) for a discussion on the subject.

In addition to altering the default rendering, excerpts have additional rendering options as shown below.

7.3. Excerpting commonly accessed data

A common situation with REST services arises when you compose domain objects. For example, a `Person` is stored in one table and their related `Address` is stored in another. By default, Spring Data REST will serve up the person's `address` as a URI the client must navigate. But if it's common for consumers to always fetch this extra piece of data, an excerpt projection can go ahead and inline this extra piece of data, saving you an extra `GET`. To do so, let's define another excerpt projection:

```
@Projection(name = "inlineAddress", types = { Person.class }) ①
interface InlineAddress {

    String getFirstName();

    String getLastName();

    Address getAddress(); ②
}
```

① This projection has been named `inlineAddress`.

② This projection adds in `getAddress` which returns the `Address` field. When used inside a projection, it causes the information to be inlined.

We can plug it into the `PersonRepository` definition as follows:

```
@RepositoryRestResource(excerptProjection = InlineAddress.class)
interface PersonRepository extends CrudRepository<Person, Long> {}
```

This will cause the HAL document to appear as follows:

```
{
  "firstName" : "Frodo",
  "lastName" : "Baggins",
  "address" : { ①
    "street": "Bag End",
    "state": "The Shire",
    "country": "Middle Earth"
  },
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons/1"
    },
    "address" : { ②
      "href" : "http://localhost:8080/persons/1/address"
    }
  }
}
```

This should appear as a mix of what you've seen so far.

- ① The `address` data is inlined directly, so you don't have to navigate to get it.
- ② The link to the `Address` resource is still provided, making it still possible to navigate to its own resource.

WARNING

Configuring `@RepositoryRestResource(excerptProjection=...)` for a repository alters the default behavior. This can potentially cause a breaking change to consumers of your service if you have already made a release. Use with caution.

Chapter 8. Conditional Operations with Headers

This section shows how Spring Data REST uses standard HTTP headers to enhance performance, conditionalize operations, and easily contribute to a more sophisticated frontend.

8.1. ETag, If-Match, and If-None-Match headers

The [ETag header](#) provides a way to tag resources. This can prevent clients from overriding each other while also making it possible to reduce unnecessary calls.

Example 3. A POJO with a version number

```
public class Sample {  
  
    @Version Long version; ①  
  
    Sample(Long version) {  
        this.version = version;  
    }  
}
```

① Spring Data Commons's [@Version](#) annotation flags this field as a version marker.

This POJO, when served up as a REST resource by Spring Data REST, will have an [ETag](#) header with the value of the version field.

We can conditionally [PUT](#), [PATCH](#), or [DELETE](#) that resource if we supply a [If-Match](#) header like this:

```
curl -v -X PATCH -H 'If-Match: <value of previous ETag>' ...
```

Only if the resource's current ETag state matches this [If-Match](#) header will the operation be carried out. This prevents clients from stomping on each other. Two different clients can fetch the resource and have an identical ETag. If one client updates the resource, it will get a new ETag in the response. But the first client still has the old header. If that client attempts an update with the [If-Match](#) header, the update will fail because they no longer match. Instead, that client will receive an [HTTP 412 Precondition Failed](#) message to be sent back. The client can then catch up however is necessary.

WARNING

The term "version" may carry different semantics with different data stores, and even different semantics within your application. Spring Data REST effectively delegates to the data store's metamodel to discern if a field is versioned, and if so, only allow the listed updates if ETags match.

The **If-None-Match header** provides an alternative. Instead of conditional updates, **If-None-Match** allow conditional queries.

```
curl -v -H 'If-None-Match: <value of previous etag>' ...
```

This command (by default) executes a **GET**. Spring Data REST will check for **If-None-Match** headers while doing a GET. If the header **MATCHES** the ETag, it will conclude nothing has changed, and instead of sending a copy of the resource, instead send back an HTTP **304 Not Modified** status code. Semantically, it reads "If this supplied header value doesn't match the server-side version, then send me the whole resource, otherwise, don't send me anything."

NOTE

This POJO is from an ETag-based unit test, so it doesn't have **@Entity** (JPA) or **@Document** (MongoDB) annotations as expected in application code. It simply focuses on how a field with **@Version** results in an ETag header.

8.2. If-Modified-Since header

The **If-Modified-Since header** provides a way to check if a resource has been updated since the last request to avoid resending the same data.

Example 4. The last modification date captured in a domain type

```
@Document
public class Receipt {

    public @Id String id;
    public @Version Long version;
    public @LastModifiedDate Date date; ①

    public String saleItem;
    public BigDecimal amount;

}
```

- ① Spring Data Commons's **@LastModifiedDate** annotation allows capturing this information in multiple formats (JodaTime's **DateTime**, legacy Java **Date** and **Calendar**, JDK8 date/time types, as well as **long/Long**).

With this field, Spring Data REST will return a **Last-Modified** header like this:

```
Last-Modified: Wed, 24 Jun 2015 20:28:15 GMT
```

This value can be capture and used for subsequent queries to avoid fetching the same data twice if it hasn't been updated.

```
curl -H "If-Modified-Since: Wed, 24 Jun 2015 20:28:15 GMT" ...
```

With this simple command, you are asking that a resource only be fetched if it's changed since that time. If so, you will get a revised **Last-Modified** header to update the client. If not, you will receive an HTTP **304 Not Modified** status code.

The header is perfectly formatted to send back for a future query.

WARNING

Don't mix and match header value with different queries. Results could be disastrous. ONLY use the header values when you are requesting the exact same URI and parameters.

8.3. Architecting a more efficient frontend

ETags combined with the **If-Match** and **If-None-Match** headers empower you to build frontend that is more friendly to consumer's data plans and mobile battery lives.

1. Identify the entities that need locking and add a version attribute. HTML5 nicely supports data-* attributes, so store it in the DOM somewhere like **data-etag** attribute.
2. Identify the entities that would benefit from tracking most recent updates. When fetching these resources, store the **Last-Modified** in the DOM (**data-last-modified** perhaps).
3. When fetching resources, also embed self URIs into your DOM nodes (perhaps **data-uri** or **data-self**) so it's very easy to go back to the resource.
4. Adjust **PUT/PATCH/DELETE** operations to use **If-Match** and also handle HTTP **412 Precondition Failed** status codes.
5. Adjust **GET** operations to use **If-None-Match**, **If-Modified-Since**, and also handle HTTP **304 Not Modified** status codes.

By embedding **ETags** and **Last-Modified** values into your DOM (or perhaps elsewhere for a native mobile app), you can then reduce the consumption of data/battery by NOT retrieving the same thing over and over. You can also avoid colliding with other clients, and instead, be alerted when you need to reconcile differences.

Hence, with just a little tweaking on your frontend and some entity-level edits, the backend will serve up time sensitive details you can cash in on when building a customer-friendly client.

Chapter 9. Validation

There are two ways to register a `Validator` instance in Spring Data REST: wire it by bean name or register the validator manually. For the majority of cases, the simple bean name prefix style will be sufficient.

In order to tell Spring Data REST you want a particular `Validator` assigned to a particular event, you simply prefix the bean name with the event you're interested in. For example, to validate instances of the `Person` class before new ones are saved into the repository, you would declare an instance of a `Validator<Person>` in your `ApplicationContext` with the bean name "beforeCreatePersonValidator". Since the prefix "beforeCreate" matches a known Spring Data REST event, that validator will be wired to the correct event.

9.1. Assigning Validators manually

If you would rather not use the bean name prefix approach, then you simply need to register an instance of your validator with the bean whose job it is to invoke validators after the correct event. In your configuration that implements `RepositoryRestConfigurer` or subclasses Spring Data REST's `RepositoryRestConfigurerAdapter`, override the `configureValidatingRepositoryEventListener` method and call `addValidator` on the `ValidatingRepositoryEventListener`, passing the event you want this validator to be triggered on, and an instance of the validator.

```
@Override
protected void configureValidatingRepositoryEventListener
(ValidatingRepositoryEventListener v) {
    v.addValidator("beforeSave", new BeforeSaveValidator());
}
```

Chapter 10. Events

There are eight different events that the REST exporter emits throughout the process of working with an entity. Those are:

- `BeforeCreateEvent`
- `AfterCreateEvent`
- `BeforeSaveEvent`
- `AfterSaveEvent`
- `BeforeLinkSaveEvent`
- `AfterLinkSaveEvent`
- `BeforeDeleteEvent`
- `AfterDeleteEvent`

10.1. Writing an ApplicationListener

There is an abstract class you can subclass which listens for these kinds of events and calls the appropriate method based on the event type. You just override the methods for the events you're interested in.

```
public class BeforeSaveEventListener extends AbstractRepositoryEventListener {

    @Override
    public void onBeforeSave(Object entity) {
        ... logic to handle inspecting the entity before the Repository saves it
    }

    @Override
    public void onAfterDelete(Object entity) {
        ... send a message that this entity has been deleted
    }
}
```

One thing to note with this approach, however, is that it makes no distinction based on the type of the entity. You'll have to inspect that yourself.

10.2. Writing an annotated handler

Another approach is to use an annotated handler, which does filter events based on domain type.

To declare a handler, create a POJO and put the `@RepositoryEventHandler` annotation on it. This tells the `BeanPostProcessor` that this class needs to be inspected for handler methods.

Once it finds a bean with this annotation, it iterates over the exposed methods and looks for annotations that correspond to the event you're interested in. For example, to handle `BeforeSaveEvent`'s in an annotated POJO for different kinds of domain types, you'd define your

class like this:

```
@RepositoryEventHandler ①
public class PersonEventHandler {

    @HandleBeforeSave
    public void handlePersonSave(Person p) {
        // ... you can now deal with Person in a type-safe way
    }

    @HandleBeforeSave
    public void handleProfileSave(Profile p) {
        // ... you can now deal with Profile in a type-safe way
    }
}
```

① It's possible to narrow the types this handler applies against by using `@RepositoryEventHandler(Person.class)`.

The domain type whose events you're interested in is determined from the type of the first parameter of the annotated methods.

To register your event handler, either mark the class with one of Spring's `@Component` stereotypes so it can be picked up by `@SpringBootApplication` or `@ComponentScan`. Or declare an instance of your annotated bean in your `ApplicationContext`. Then the `BeanPostProcessor` that is created in `RepositoryRestMvcConfiguration` will inspect the bean for handlers and wire them to the correct events.

```
@Configuration
public class RepositoryConfiguration {

    @Bean
    PersonEventHandler personEventHandler() {
        return new PersonEventHandler();
    }
}
```

NOTE

Spring Data REST events are customized [Spring application events](#). Spring events are synchronous by default, unless they get republished across a boundary (like issuing a WebSocket event or crossing into a thread).

Chapter 11. Integration

This section details various ways to integrate with Spring Data REST components, whether from a Spring application that is using Spring Data REST or from other means.

11.1. Programmatic Links

Sometimes you need to add links to exported resources in your own custom built Spring MVC controllers. There are three basic levels of linking available:

1. Manually assembling links
2. Using Spring HATEOAS's [LinkBuilder](#) with `linkTo()`, `slash()`, etc.
3. Using Spring Data REST's implementation of [RepositoryEntityLinks](#).

The first suggestion is terrible and should be avoided at all costs. It makes your code brittle and high risk. The second is handy when creating links to other hand written Spring MVC controllers. The last one, which you'll see in a moment, is good for looking up resource links that are exported by Spring Data REST.

Assuming you have configured your code to use Spring's autowiring,

```
public class MyWebApp {  
  
    private RepositoryEntityLinks entityLinks;  
  
    @Autowired  
    public MyWebApp(RepositoryEntityLinks entityLinks) {  
        this.entityLinks = entityLinks;  
    }  
}
```

...you can then use the following operations:

Table 3. Ways to link to exported resources

Method	Description
<code>entityLinks.linkToCollectionResource(Person.class)</code>	Provide a link to the collection resource of that type.
<code>entityLinks.linkToSingleResource(Person.class, 1)</code>	Provide a link to a single resource.
<code>entityLinks.linkToPagedResource(Person.class, new PageRequest(...))</code>	Provide a link to a paged resource.
<code>entityLinks.linksToSearchResources(Person.class)</code>	Provides a list of links for all the finder methods exposed by the corresponding repository.
<code>entityLinks.linkToSearchResource(Person.class, "findByLastName")</code>	Provide a finder link by rel, i.e. the name of the finder.

NOTE

All of the search-based links support extra parameters for paging and sorting. Checkout [RepositoryEntityLinks](#) for specifics. There is also `linkFor(Class<?> type)`, but that returns a Spring HATEOAS `LinkBuilder`, which returns you to the lower level API. Try to use the other ones first.

Chapter 12. Metadata

This section details the various forms of metadata provided by a Spring Data REST-based application.

12.1. Application-Level Profile Semantics (ALPS)

ALPS is a data format for defining simple descriptions of application-level semantics, similar in complexity to HTML microformats. An ALPS document can be used as a profile to explain the application semantics of a document with an application-agnostic media type (such as HTML, HAL, Collection+JSON, Siren, etc.). This increases the reusability of profile documents across media types.

— M. Amundsen / L. Richardson / M. Foster, <http://tools.ietf.org/html/draft-amundsen-richardson-foster-alps-00>

Spring Data REST provides an ALPS document for every exported repository. It contains information about both the RESTful transitions as well as the attributes of each repository.

At the root of a Spring Data REST app is a **profile** link. Assuming you had an app with both **persons** and related **addresses**, the root document would look like this:

```
{
  "_links" : {
    "persons" : {
      "href" : "http://localhost:8080/persons"
    },
    "addresses" : {
      "href" : "http://localhost:8080/addresses"
    },
    "profile" : {
      "href" : "http://localhost:8080/profile"
    }
  }
}
```

A **profile** link, as defined in [RFC 6906](#), is a place to include application level details. The [ALPS draft spec](#) is meant to define a particular profile format which we'll explore further down in this section.

If you navigate into the **profile** link at localhost:8080/profile, you would see something like this:

```
{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/profile"
    },
    "persons" : {
      "href" : "http://localhost:8080/profile/persons"
    },
    "addresses" : {
      "href" : "http://localhost:8080/profile/addresses"
    }
  }
}
```

IMPORTANT

At the root level, **profile** is a single link and hence can't handle serving up more than one application profile. That is why you must navigate to **/profile** to find a link for each resource's metadata.

Let's navigate to **/profile/persons** and look at the profile data for a **Person** resource.

```

{
  "version" : "1.0",
  "descriptors" : [ {
    "id" : "person-representation", ①
    "descriptors" : [ {
      "name" : "firstName",
      "type" : "SEMANTIC"
    }, {
      "name" : "lastName",
      "type" : "SEMANTIC"
    }, {
      "name" : "id",
      "type" : "SEMANTIC"
    }, {
      "name" : "address",
      "type" : "SAFE",
      "rt" : "http://localhost:8080/profile/addresses#address"
    } ]
  }, {
    "id" : "create-persons", ②
    "name" : "persons", ③
    "type" : "UNSAFE", ④
    "rt" : "#person-representation" ⑤
  }, {
    "id" : "get-persons",
    "name" : "persons",
    "type" : "SAFE",
    "rt" : "#person-representation"
  }, {
    "id" : "delete-person",
    "name" : "person",
    "type" : "IDEMPOTENT",
    "rt" : "#person-representation"
  }, {
    "id" : "patch-person",
    "name" : "person",
    "type" : "UNSAFE",
    "rt" : "#person-representation"
  }, {
    "id" : "update-person",
    "name" : "person",
    "type" : "IDEMPOTENT",
    "rt" : "#person-representation"
  }, {
    "id" : "get-person",
    "name" : "person",
    "type" : "SAFE",
    "rt" : "#person-representation"
  } ]
}

```


- ① At the top is a detailed listing of the attributes of a **Person** resource, identified as **#person-representation**. It lists the names of the attributes.
- ② After the resource representation are all the supported operations. This one is how to create a new **Person**.
- ③ The name is **persons**, which indicates that a POST should be applied to the whole collection, not a single **person**.
- ④ The **type** is **UNSAFE** because this operation can alter the state of the system.

NOTE

This JSON document has a media type of **application/alps+json**. This is different than the previous JSON document, which had a media type of **application/hal+json**. These formats are different and governed by different specs.

You will also find a "profile" link shown in the collection of **_links** when you are looking at a collection resource.

```
{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons" ①
    },
    ... other links ...
    "profile" : {
      "href" : "http://localhost:8080/profile/persons" ②
    }
  },
  ...
}
```

- ① This HAL document represents the **Person** collection.
- ② It has a **profile** link to the same URI for metadata.

The **profile** link, again, will serve up ALPS by default or if you use an **Accept header** of **application/alps+json**.

12.1.1. Hypermedia control types

ALPS displays types for each hypermedia control. They include:

Table 4. ALPS types

Type	Description
SEMANTIC	A state element (e.g. HTML.SPAN, HTML.INPUT, etc.).
SAFE	A hypermedia control that triggers a safe, idempotent state transition (e.g. GET or HEAD).
IDEMPOTENT	A hypermedia control that triggers an unsafe, idempotent state transition (e.g. PUT or DELETE).

Type	Description
UNSAFE	A hypermedia control that triggers an unsafe, non-idempotent state transition (e.g. POST).

In the representation section up above, bits of data from the application are marked **SEMANTIC**. The **address** field is a link that involves a safe **GET** to retrieve. Hence, it is marked **SAFE**. Hypermedia operations themselves map onto the types as shown the table.

12.1.2. ALPS with Projections

If you define any projections, they are also listed in the ALPS metadata. Assuming we also defined **inlineAddress** and **noAddresses**, they would appear inside the relevant operations, i.e. **GET** for the whole collection as well **GET** for a single resource. The following shows the alternate version of the **get-persons** subsection:

```

...
{
  "id" : "get-persons",
  "name" : "persons",
  "type" : "SAFE",
  "rt" : "#person-representation",
  "descriptors" : [ { ①
    "name" : "projection",
    "doc" : {
      "value" : "The projection that shall be applied when rendering the response.
Acceptable values available in nested descriptors.",
      "format" : "TEXT"
    },
    "type" : "SEMANTIC",
    "descriptors" : [ {
      "name" : "inlineAddress", ②
      "type" : "SEMANTIC",
      "descriptors" : [ {
        "name" : "address",
        "type" : "SEMANTIC"
      }, {
        "name" : "firstName",
        "type" : "SEMANTIC"
      }, {
        "name" : "lastName",
        "type" : "SEMANTIC"
      } ]
    }, {
      "name" : "noAddresses", ③
      "type" : "SEMANTIC",
      "descriptors" : [ {
        "name" : "firstName",
        "type" : "SEMANTIC"
      }, {
        "name" : "lastName",
        "type" : "SEMANTIC"
      } ]
    } ]
  } ]
}
...

```

- ① A new attribute, **descriptors**, appears containing an array with one entry, **projection**.
- ② Inside the **projection.descriptors** we can see **inLineAddress** listed. It will render **address**, **firstName**, and **lastName**. Relationships rendered inside a projection result in inlining the data fields.
- ③ Also found is **noAddresses**, which serves up a subset containing **firstName** and **lastName**.

With all this information, a client should be able to deduce not only the RESTful transitions

available, but also, to some degree, the data elements needed to interact.

12.1.3. Adding custom details to your ALPS descriptions

It's possible to create custom messages that appear in your ALPS metadata. Just create `rest-messages.properties` like this:

```
rest.description.person=A collection of people
rest.description.person.id=primary key used internally to store a person (not for
RESTful usage)
rest.description.person.firstName=Person's first name
rest.description.person.lastName=Person's last name
rest.description.person.address=Person's address
```

As you can see, this defines details to display for a `Person` resource. They alter the ALPS format of the **person-representation** as follows:

```

...
{
  "id" : "person-representation",
  "doc" : {
    "value" : "A collection of people", ①
    "format" : "TEXT"
  },
  "descriptors" : [ {
    "name" : "firstName",
    "doc" : {
      "value" : "Person's first name", ②
      "format" : "TEXT"
    },
    "type" : "SEMANTIC"
  }, {
    "name" : "lastName",
    "doc" : {
      "value" : "Person's last name", ③
      "format" : "TEXT"
    },
    "type" : "SEMANTIC"
  }, {
    "name" : "id",
    "doc" : {
      "value" : "primary key used internally to store a person (not for RESTful
usage)", ④
      "format" : "TEXT"
    },
    "type" : "SEMANTIC"
  }, {
    "name" : "address",
    "doc" : {
      "value" : "Person's address", ⑤
      "format" : "TEXT"
    },
    "type" : "SAFE",
    "rt" : "http://localhost:8080/profile/addresses#address"
  } ]
}
...

```

By supplying these property settings, each field has an extra **doc** attribute.

- ① The value of `rest.description.person` maps into the whole representation.
- ② The value of `rest.description.person.firstName` maps to the **firstName** attribute.
- ③ The value of `rest.description.person.lastName` maps to the **lastName** attribute.
- ④ The value of `rest.description.person.id` maps to the **id** attribute, a field not normally displayed.
- ⑤ The value of `rest.description.person.address` maps to the **address** attribute.

NOTE

Spring MVC (which is the essence of a Spring Data REST application) supports locales, meaning you can bundle up multiple properties files with different messages.

12.2. JSON Schema

[JSON Schema](#) is another form of metadata supported by Spring Data REST. Per their website, JSON Schema has the following advantages:

- describes your existing data format
- clear, human- and machine-readable documentation
- complete structural validation, useful for automated testing and validating client-submitted data

As shown in the [previous section](#), you can reach this data by navigating from the root URI to the "profile" link.

```
{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/profile"
    },
    "persons" : {
      "href" : "http://localhost:8080/profile/persons"
    },
    "addresses" : {
      "href" : "http://localhost:8080/profile/addresses"
    }
  }
}
```

These links are the same as shown earlier. To retrieve JSON Schema you invoke them with Accept header **application/schema+json**.

In this case, if you executed `curl -H 'Accept:application/schema+json' http://localhost:8080/profile/persons`, you would see something like this:

```

{
  "title" : "org.springframework.data.rest.webmvc.jpa.Person", ①
  "properties" : { ②
    "firstName" : {
      "readOnly" : false,
      "type" : "string"
    },
    "lastName" : {
      "readOnly" : false,
      "type" : "string"
    },
    "siblings" : {
      "readOnly" : false,
      "type" : "string",
      "format" : "uri"
    },
    "created" : {
      "readOnly" : false,
      "type" : "string",
      "format" : "date-time"
    },
    "father" : {
      "readOnly" : false,
      "type" : "string",
      "format" : "uri"
    },
    "weight" : {
      "readOnly" : false,
      "type" : "integer"
    },
    "height" : {
      "readOnly" : false,
      "type" : "integer"
    }
  },
  "descriptors" : { },
  "type" : "object",
  "$schema" : "http://json-schema.org/draft-04/schema#"
}

```

① The type that was exported

② A listing of properties

There are more details if your resources have links to other resources.

You will also find a "profile" link shown in the collection of **_links** when you are looking at a collection resource.

```
{
  "_links" : {
    "self" : {
      "href" : "http://localhost:8080/persons" ①
    },
    ... other links ...
    "profile" : {
      "href" : "http://localhost:8080/profile/persons" ②
    }
  },
  ...
}
```

① This HAL document represents the **Person** collection.

② It has a **profile** link to the same URI for metadata.

The **profile** link, again, will serve up **ALPS** by default. If you supply it with an **Accept header** of **application/schema+json**, it will render the JSON Schema representation.

Chapter 13. Security

Spring Data REST works quite well with Spring Security. This section will show examples of how to secure your Spring Data REST services with **method level security**.

13.1. @Pre and @Post security

The following example from Spring Data REST's test suite shows Spring Security's [PreAuthorization model](#), the most sophisticated version:

Example 5. `spring-data-rest-tests/spring-data-rest-tests-security/src/test/java/org/springframework/data/rest/tests/security/PreAuthorizedOrderRepository.java`

```
@PreAuthorize("hasRole('ROLE_USER')") ①
public interface PreAuthorizedOrderRepository extends CrudRepository<Order, UUID>
{

    @PreAuthorize("hasRole('ROLE_ADMIN')") ②
    @Override
    void delete(UUID aLong);

    @PreAuthorize("hasRole('ROLE_ADMIN')")
    @Override
    void delete(Order order);

    @PreAuthorize("hasRole('ROLE_ADMIN')")
    @Override
    void delete(Iterable<? extends Order> orders);

    @PreAuthorize("hasRole('ROLE_ADMIN')")
    @Override
    void deleteAll();
}
```

This is a standard Spring Data repository definition extending `CrudRepository` with some key changes.

- ① This Spring Security annotation secures the entire repository. The [Spring Security SpEL expression](#) indicates the principal must have **ROLE_USER** in his collection of roles.
- ② To change method-level settings, you must override the method signature and apply a Spring Security annotation. In this case, the method overrides the repository-level settings with the requirement that the user have **ROLE_ADMIN** to perform a delete.

IMPORTANT

Repository and method level security settings don't combine. Instead, method-level settings *override* repository level settings.

The previous example illustrates that `CrudRepository`, in fact, has four delete methods. You must override all delete methods to properly secure it.

13.2. @Secured security

The following example shows Spring Security's older `@Secured` annotation, which is purely role-based:

Example 6. spring-data-rest-tests/spring-data-rest-tests-

security/src/test/java/org/springframework/data/rest/tests/security/SecuredPersonRepository.java

```
@Secured("ROLE_USER") ①
@RepositoryRestResource(collectionResourceRel = "people", path = "people")
public interface SecuredPersonRepository extends CrudRepository<Person, UUID> {

    @Secured("ROLE_ADMIN") ②
    @Override
    void delete(UUID aLong);

    @Secured("ROLE_ADMIN")
    @Override
    void delete(Person person);

    @Secured("ROLE_ADMIN")
    @Override
    void delete(Iterable<? extends Person> persons);

    @Secured("ROLE_ADMIN")
    @Override
    void deleteAll();
}
```

- ① This results in the same security check, but has less flexibility. It only allows roles as the means to restrict access.
- ② Again, this shows that delete methods require **ROLE_ADMIN**.

NOTE

If you are starting with a new project or first applying Spring Security, `@PreAuthorize` is the recommended solution. If are already using Spring Security with `@Secured` in other parts of your app, you can continue on that path without rewriting everything.

13.3. Enabling method level security

To configure method level security, here is a brief snippet from Spring Data REST's test suite:

Example 7. *spring-data-rest-tests/spring-data-rest-tests-security/src/test/java/org/springframework/data/rest/tests/security/SecurityConfiguration.java*

```
@Configuration ①
@EnableWebSecurity
@EnableGlobalMethodSecurity(securedEnabled = true, prePostEnabled = true) ②
public class SecurityConfiguration extends WebSecurityConfigurerAdapter { ③
    ...
}
```

- ① This is a Spring configuration class.
- ② It uses Spring Security's `@EnableGlobalMethodSecurity` annotation to enable both `@Secured` and `@Pre/@Post` support. NOTE: You don't have to use both. This particular case is used to prove both versions work with Spring Data REST.
- ③ This class extends Spring Security's `WebSecurityConfigurerAdapter` which is used for pure Java configuration of security.

The rest of the configuration class isn't listed because it follows [standard practices](#) you can read about in the Spring Security reference docs.

Chapter 14. Tools

14.1. The HAL Browser

The developer of the [HAL spec](#) has a useful application: [the HAL Browser](#). It's a web app that stirs in a little HAL-powered JavaScript. You can point it at any Spring Data REST API and use it to navigate the app and create new resources.

Instead of pulling down the files, embedding them in your application, and crafting a Spring MVC controller to serve them up, all you need to do is add a single dependency.

In Maven:

```
<dependencies>
  <dependency>
    <groupId>org.springframework.data</groupId>
    <artifactId>spring-data-rest-hal-browser</artifactId>
  </dependency>
</dependencies>
```

In Gradle:

```
dependencies {
    compile 'org.springframework.data:spring-data-rest-hal-browser'
}
```

NOTE

If you use Spring Boot or the Spring Data BOM (bill of materials), you don't need to specify the version.

This dependency will autoconfigure the HAL Browser to be served up when you visit your application's root URI in a browser. (NOTE: <http://localhost:8080> was plugged into the browser, and it redirect to the URL shown below.)

The screenshot shows a web browser window titled "The HAL Browser (customi: x)" with the address bar displaying "localhost:8080/browser/index.html#/" and a "Go!" button. The application interface is divided into two main sections: "Explorer" on the left and "Inspector" on the right.

Explorer Section:

- Custom Request Headers:** A text input field with a "Go!" button.
- Properties:** A text input field containing "{}".
- Links:** A table listing API resources with columns for "rel", "title", "name / index", "docs", "GET", and "NON-GET".

rel	title	name / index	docs	GET	NON-GET
persons					
addresses					
profile					

Inspector Section:

- Response Headers:** Displays the status "200 OK" and headers: "Date: Mon, 03 Aug 2015 19:55:45 GMT", "Server: Apache-Coyote/1.1", "Transfer-Encoding: chunked", and "Content-Type: application/hal+json; charset=UTF-8".
- Response Body:** Displays a JSON HAL document:


```
{
  "_links": {
    "persons": {
      "href": "http://localhost:8080/persons{?projection}",
      "templated": true
    },
    "addresses": {
      "href": "http://localhost:8080/addresses"
    },
    "profile": {
      "href": "http://localhost:8080/alps"
    }
  }
}
```

The screen shot above shows the root path of the API. On the right side are details from the response including headers and the body (a HAL document).

The HAL Browser reads the links from the response and puts them on a list on the left side. You can either click on the **GET** button and navigate to one of the collections, or click on the **non-GET** option to make changes.

The HAL Browser speaks **URI Template**. You may notice up above the **GET** button next to **persons** has a question mark icon. An expansion dialog will pop-up if you choose to navigate to it like this:

Expand URI Template



URI Template:

```
http://localhost:8080/persons{?projection}
```

Input (JSON):

```
{  
  "projection": ""  
}
```

Expanded URI:

```
http://localhost:8080/persons
```

Follow URI

If you click **Follow URI** without entering anything, the variables will essentially be ignored. For situations like [Projections and Excerpts](#) or [Paging and Sorting](#), this can be useful.

When you click on a **non-GET** button, a pop-up dialog appears. By default, it shows **POST**. This field can be adjusted to either **PUT** or **PATCH**. The headers are filled out to properly to submit a new JSON document.

Below the URI, method, and headers are the fields. These are automatically supplied based on the metadata of the resources, automatically generated by Spring Data REST. Update your domain objects, and the pop-up will reflect it.

Person

First name

Last name

Action:

POST

http://localhost:8080/persons

Make Request

Chapter 15. Customizing Spring Data REST

There are many options to tailor Spring Data REST. These subsections show how.

15.1. Customizing item resource URIs

By default the URI for item resources are comprised of the path segment used for the collection resource with the database identifier appended. That allows us to use the repository's `findOne(...)` method to lookup entity instances. As of Spring Data REST 2.5 this can be customized by using configuration API on `RepositoryRestConfiguration` (preferred on Java 8) or by registering an implementation of `EntityLookup` as Spring bean in your application. Spring Data REST will pick those up and tweak the URI generation according to their implementation.

Assume a `User` with a `username` property that uniquely identifies it. Also, assume we have a method `Optional<User> findByUsername(String username)` on the according repository.

On Java 8 we can simply register the mapping methods as method references to weak the URI creation as follows:

```
@Component
public class SpringDataRestCustomization extends RepositoryRestConfigurerAdapter {

    @Override
    public void configureRepositoryRestConfiguration(RepositoryRestConfiguration config)
    {

        config.withCustomEntityLookup().//
            forRepository(UserRepository.class, User::getUsername, UserRepository:
:findByUsername);
    }
}
```

`forRepository(...)` takes the repository type as first argument a method reference mapping the repositories domain type to some target type, as well as another method reference to map that value back using the repository mentioned as first argument.

If you're not running Java 8 or better, you could use the method but it would require a few quite verbose anonymous inner classes to be use. That's why on older Java versions you probably prefer implementing a `UserEntityLookup` looking like this:


```

@Component
public class UserEntityLookup extends EntityLookupSupport<User> {

    private final UserRepository repository;

    public UserEntityLookup(UserRepository repository) {
        this.repository = repository;
    }

    @Override
    public Serializable getResourceIdentifier(User entity) {
        return entity.getUsername();
    }

    @Override
    public Object lookupEntity(Serializable id) {
        return repository.findByUsername(id.toString());
    }
}

```

Note, how `getResourceIdentifier(...)` returns the username to be used by the URI creation. To load entity instances by the value returned from that method we now implement `lookupEntity(...)` using the query method available on the `UserRepository`.

15.2. Configuring the REST URL path

Configuring the segments of the URL path under which the resources of a JPA repository are exported is simple. You just add an annotation at the class level and/or at the query method level.

By default, the exporter will expose your `CrudRepository` using the name of the domain class. Spring Data REST also applies the `Evo Inflector` to pluralize this word. So a repository defined as follows:

```

interface PersonRepository extends CrudRepository<Person, Long> {}

```

Will, by default, be exposed under the URL <http://localhost:8080/persons/>

To change how the repository is exported, add a `@RestResource` annotation at the class level:

```

@RepositoryRestResource(path = "people")
interface PersonRepository extends CrudRepository<Person, Long> {}

```

Now the repository will be accessible under the URL: <http://localhost:8080/people/>

If you have query methods defined, those also default to be exposed by their name:

```
interface PersonRepository extends CrudRepository<Person, Long> {

    List<Person> findByName(String name);
}
```

This would be exposed under the URL: <http://localhost:8080/persons/search/findByName>

NOTE All query method resources are exposed under the resource `search`.

To change the segment of the URL under which this query method is exposed, use the `@RestResource` annotation again:

```
@RepositoryRestResource(path = "people")
interface PersonRepository extends CrudRepository<Person, Long> {

    @RestResource(path = "names")
    List<Person> findByName(String name);
}
```

Now this query method will be exposed under the URL: <http://localhost:8080/people/search/names>

15.2.1. Handling rels

Since these resources are all discoverable, you can also affect how the "rel" attribute is displayed in the links sent out by the exporter.

For instance, in the default configuration, if you issue a request to <http://localhost:8080/persons/search> to find out what query methods are exposed, you'll get back a list of links:

```
{
  "_links" : {
    "findByName" : {
      "href" : "http://localhost:8080/persons/search/findByName"
    }
  }
}
```

To change the rel value, use the `rel` property on the `@RestResource` annotation:

```
@RepositoryRestResource(path = "people")
interface PersonRepository extends CrudRepository<Person, Long> {

    @RestResource(path = "names", rel = "names")
    List<Person> findByName(String name);
}
```

This would result in a link value of:

```
{
  "_links" : {
    "names" : {
      "href" : "http://localhost:8080/persons/search/names"
    }
  }
}
```

NOTE

These snippets of JSON assume you are using Spring Data REST's default format of [HAL](#). It's possible to turn off HAL, which would cause the output to look different. But your ability to override rel names is totally independent of the rendering format.

```
@RepositoryRestResource(path = "people", rel = "people")
interface PersonRepository extends CrudRepository<Person, Long> {

    @RestResource(path = "names", rel = "names")
    List<Person> findByName(String name);
}
```

Altering the rel of a Repository changes the top level name:

```
{
  "_links" : {
    "people" : {
      "href" : "http://localhost:8080/people"
    },
    ...
  }
}
```

In the top level fragment above:

- `path = "people"` changed the value in `href` from `/persons` to `/people`
- `rel = "people"` changed the name of that link from `persons` to `people`

When you navigate to the **search** resource of this repository, the finder-method's `@RestResource` annotation has altered the path as shown below:

```
{
  "_links" : {
    "names" : {
      "href" : "http://localhost:8080/people/search/names"
    }
  }
}
```

This collection of annotations in your Repository definition has caused the following changes:

- The Repository-level annotation's `path = "people"` is reflected in the base URI with `/people`
- Being a finder method provides you with `/people/search`
- `path = "names"` creates a URI of `/people/search/names`
- `rel = "names"` changes the name of that link from `findByName` to `names`

15.2.2. Hiding certain repositories, query methods, or fields

You may not want a certain repository, a query method on a repository, or a field of your entity to be exported at all. Examples include hiding fields like `password` on a `User` object or similar sensitive data. To tell the exporter to not export these items, annotate them with `@RestResource` and set `exported = false`.

For example, to skip exporting a Repository:

```
@RepositoryRestResource(exported = false)
interface PersonRepository extends CrudRepository<Person, Long> {}
```

To skip exporting a query method:

```
@RepositoryRestResource(path = "people", rel = "people")
interface PersonRepository extends CrudRepository<Person, Long> {

    @RestResource(exported = false)
    List<Person> findByName(String name);
}
```

Or to skip exporting a field:

```

@Entity
public class Person {

    @Id @GeneratedValue private Long id;

    @OneToMany
    @RestResource(exported = false)
    private Map<String, Profile> profiles;
}

```

WARNING

Projections provide the means to change what is exported and effectively [side step these settings](#). If you create any projections against the same domain object, it's your responsibility to NOT export the fields. See

15.2.3. Hiding repository CRUD methods

If you don't want to expose a save or delete method on your `CrudRepository`, you can use the `@RestResource(exported = false)` setting by overriding the method you want to turn off and placing the annotation on the overridden version. For example, to prevent HTTP users from invoking the delete methods of `CrudRepository`, override all of them and add the annotation to the overridden methods.

```

@RepositoryRestResource(path = "people", rel = "people")
interface PersonRepository extends CrudRepository<Person, Long> {

    @Override
    @RestResource(exported = false)
    void delete(Long id);

    @Override
    @RestResource(exported = false)
    void delete(Person entity);
}

```

WARNING

It is important that you override *both* delete methods as the exporter currently uses a somewhat naive algorithm for determining which CRUD method to use in the interest of faster runtime performance. It's not currently possible to turn off the version of delete which takes an ID but leave exported the version that takes an entity instance. For the time being, you can either export the delete methods or not. If you want turn them off, then just keep in mind you have to annotate both versions with `exported = false`.

15.3. Adding Spring Data REST to an existing Spring MVC Application

NOTE

The following steps are unnecessary if you are using Spring Boot. Adding **spring-boot-starter-data-rest** will cause it to automatically get added to your application.

If you have an existing Spring MVC application and you'd like to integrate Spring Data REST, it's actually very easy.

Somewhere in your Spring MVC configuration (most likely where you configure your MVC resources) add a bean reference to the JavaConfig class that is responsible for configuring the `RepositoryRestController`. The class name is `org.springframework.data.rest.webmvc.RepositoryRestMvcConfiguration`.

In Java, this would look like:

```
import org.springframework.context.annotation.Import;
import org.springframework.data.rest.webmvc.RepositoryRestMvcConfiguration;

@Configuration
@Import(RepositoryRestMvcConfiguration.class)
public class MyApplicationConfiguration {

    ...
}
```

In XML this would look like:

```
<bean class=
"org.springframework.data.rest.webmvc.config.RepositoryRestMvcConfiguration"/>
```

When your `ApplicationContext` comes across this bean definition it will bootstrap the necessary Spring MVC resources to fully-configure the controller for exporting the repositories it finds in that `ApplicationContext` and any parent contexts.

15.3.1. More on required configuration

There are a couple Spring MVC resources that Spring Data REST depends on that must be configured correctly for it to work inside an existing Spring MVC application. We've tried to isolate those resources from whatever similar resources already exist within your application, but it may be that you want to customize some of the behavior of Spring Data REST by modifying these MVC components.

The most important things that we configure especially for use by Spring Data REST include:

RepositoryRestHandlerMapping

We register a custom `HandlerMapping` instance that responds only to the `RepositoryRestController` and only if a path is meant to be handled by Spring Data REST. In order to keep paths that are meant to be handled by your application separate from those handled by Spring Data REST, this custom `HandlerMapping` inspects the URL path and checks to see if a Repository has been exported under that name. If it has, it allows the request to be handled by Spring Data REST. If there is no Repository exported under that name, it returns `null`, which just means "let other `HandlerMapping` instances try to service this request".

The Spring Data REST `HandlerMapping` is configured with `order=(Ordered.LOWEST_PRECEDENCE - 100)` which means it will usually be first in line when it comes time to map a URL path. Your existing application will never get a chance to service a request that is meant for a repository. For example, if you have a repository exported under the name "person", then all requests to your application that start with `/person` will be handled by Spring Data REST and your application will never see that request. If your repository is exported under a different name, however (like "people"), then requests to `/people` will go to Spring Data REST and requests to `/person` will be handled by your application.

15.4. Overriding Spring Data REST Response Handlers

Sometimes you may want to write a custom handler for a specific resource. To take advantage of Spring Data REST's settings, message converters, exception handling, and more, use the `@RepositoryRestController` annotation instead of a standard Spring MVC `@Controller` or `@RestController`:

```

@RepositoryRestController
public class ScannerController {

    private final ScannerRepository repository;

    @Autowired
    public ScannerController(ScannerRepository repo) { ①
        repository = repo;
    }

    @RequestMapping(method = GET, value = "/scanners/search/listProducers") ②
    public @ResponseBody ResponseEntity<?> getProducers() {
        List<String> producers = repository.listProducers(); ③

        //
        // do some intermediate processing, logging, etc. with the producers
        //

        Resources<String> resources = new Resources<String>(producers); ④

        resources.add(linkTo(methodOn(ScannerController.class).getProducers())
            .withSelfRel()); ⑤

        // add other links as needed

        return ResponseEntity.ok(resources); ⑥
    }
}

```

This controller will be served from the same API base path defined in `RepositoryRestConfiguration.setBasePath` that is used by all other RESTful endpoints (e.g. `/api`). It also has these characteristics:

- ① This example uses constructor injection.
- ② This handler plugs in a custom handler for a Spring Data finder method.
- ③ This handler is using the underlying repository to fetch data, but will then do some form of post processing before returning the final data set to the client.
- ④ The results need to be wrapped up in a Spring HATEOAS `Resources` object to return a collection, but only a `Resource` for a single item.
- ⑤ Add a link back to this exact method as a "self" link.
- ⑥ Returning the collection using Spring MVC's `ResponseEntity` wrapper ensure the collection is properly wrapped and rendered in the proper accept type.

`Resources` is for a collection while `Resource` is for a single item. These types can be combined. If you know the links for each item in a collection, use `Resources<Resource<String>>` (or whatever the core domain type is). This lets you assemble links for each item as well as for the whole collection.

IMPORTANT

In this example, the combined path will be `RepositoryRestConfiguration.getBasePath() + /scanners/search/listProducers`.

If you're NOT interested in entity-specific operations but still want to build custom operations underneath `basePath`, such as Spring MVC views, resources, etc. use `@BasePathAwareController`.

WARNING

If you use `@Controller` or `@RestController` for anything, that code will be totally outside the scope of Spring Data REST. This extends to request handling, message converters, exception handling, etc.

15.5. Customizing the JSON output

Sometimes in your application you need to provide links to other resources from a particular entity. For example, a `Customer` response might be enriched with links to a current shopping cart, or links to manage resources related to that entity. Spring Data REST provides integration with [Spring HATEOAS](#) and provides an extension hook for users to alter the representation of resources going out to the client.

15.5.1. The ResourceProcessor interface

Spring HATEOAS defines a `ResourceProcessor<>` interface for processing entities. All beans of type `ResourceProcessor<Resource<T>>` will be automatically picked up by the Spring Data REST exporter and triggered when serializing an entity of type `T`.

For example, to define a processor for a `Person` entity, add a `@Bean` to your `ApplicationContext` like the following (which is taken from the Spring Data REST tests):

```
@Bean
public ResourceProcessor<Resource<Person>> personProcessor() {

    return new ResourceProcessor<Resource<Person>>() {

        @Override
        public Resource<Person> process(Resource<Person> resource) {

            resource.add(new Link("http://localhost:8080/people", "added-link"));
            return resource;
        }
    };
}
```

IMPORTANT

This example hard codes a link to `http://localhost:8080/people`. If you have a Spring MVC endpoint inside your app that you wish to link to, consider using Spring HATEOAS's `linkTo(...)` method to avoid managing the URL.

15.5.2. Adding Links

It's possible to add links to the default representation of an entity by simply calling `resource.add(Link)` like the example above. Any links you add to the `Resource` will be added to the final output.

15.5.3. Customizing the representation

The Spring Data REST exporter executes any discovered `ResourceProcessor`'s before it creates the output representation. It does this by registering a `Converter<Entity, Resource>` instance with an internal `ConversionService`. This is the component responsible for creating the links to referenced entities (e.g. those objects under the `_links` property in the object's JSON representation). It takes an `@Entity` and iterates over its properties, creating links for those properties that are managed by a `Repository` and copying across any embedded or simple properties.

If your project needs to have output in a different format, however, it's possible to completely replace the default outgoing JSON representation with your own. If you register your own `ConversionService` in the `ApplicationContext` and register your own `Converter<Person, Resource>`, then you can return a `Resource` implementation of your choosing.

15.6. Adding custom (de)serializers to Jackson's ObjectMapper

Sometimes the behavior of the Spring Data REST's `ObjectMapper`, which has been specially configured to use intelligent serializers that can turn domain objects into links and back again, may not handle your domain model correctly. There are so many ways one can structure your data that you may find your own domain model isn't being translated to JSON correctly. It's also sometimes not practical in these cases to try and support a complex domain model in a generic way. Sometimes, depending on the complexity, it's not even possible to offer a generic solution.

So to accommodate the largest percentage of the use cases, Spring Data REST tries very hard to render your object graph correctly. It will try and serialize unmanaged beans as normal POJOs and it will try and create links to managed beans where that's necessary. But if your domain model doesn't easily lend itself to reading or writing plain JSON, you may want to configure Jackson's `ObjectMapper` with your own custom type mappings and (de)serializers.

15.6.1. Abstract class registration

One key configuration point you might need to hook into is when you're using an abstract class (or an interface) in your domain model. Jackson won't know by default what implementation to create for an interface. Take the following example:

```

@Entity
public class MyEntity {

    @OneToMany
    private List<MyInterface> interfaces;
}

```

In a default configuration, Jackson has no idea what class to instantiate when POSTing new data to the exporter. This is something you'll need to tell Jackson either through an annotation, or, more cleanly, by registering a type mapping using a [Module](#).

Any [Module](#) bean declared within the scope of your [ApplicationContext](#) will be picked up by the exporter and registered with its [ObjectMapper](#). To add this special abstract class type mapping, create a [Module](#) bean and in the `setupModule` method, add an appropriate [TypeResolver](#):

```

public class MyCustomModule extends SimpleModule {

    private MyCustomModule() {
        super("MyCustomModule", new Version(1, 0, 0, "SNAPSHOT"));
    }

    @Override
    public void setupModule(SetupContext context) {
        context.addAbstractTypeResolver(
            new SimpleAbstractTypeResolver().addMapping(MyInterface.class,
                MyInterfaceImpl.class));
    }
}

```

Once you have access to the [SetupContext](#) object in your [Module](#), you can do all sorts of cool things to configure Jackson's JSON mapping. You can read more about how [Module`s work on Jackson's wiki](#).

15.6.2. Adding custom serializers for domain types

If you want to (de)serialize a domain type in a special way, you can register your own implementations with Jackson's [ObjectMapper](#) and the Spring Data REST exporter will transparently handle those domain objects correctly.

To add serializers, from your `setupModule` method implementation, do something like the following:

```

public class MyCustomModule extends SimpleModule {

    ...

    @Override
    public void setupModule(SetupContext context) {

        SimpleSerializers serializers = new SimpleSerializers();
        SimpleDeserializers deserializers = new SimpleDeserializers();

        serializers.addSerializer(MyEntity.class, new MyEntitySerializer());
        deserializers.addDeserializer(MyEntity.class, new MyEntityDeserializer());

        context.addSerializers(serializers);
        context.addDeserializers(deserializers);
    }
}

```

Now Spring Data REST will correctly handle your domain objects in case they are too complex for the 80% generic use case that Spring Data REST tries to cover.

15.7. Configuring CORS

For security reasons, browsers prohibit AJAX calls to resources residing outside the current origin. When working with client-side HTTP requests issued by a browser you want to enable specific HTTP resources to be accessible.

Spring Data REST supports as of 2.6 [Cross-origin resource sharing](#) (CORS) through [Spring's CORS](#) support.

15.7.1. Repository interface CORS configuration

You can add a `@CrossOrigin` annotation to your repository interfaces to enable CORS for the whole repository. By default `@CrossOrigin` allows all origins and HTTP methods:

```

@CrossOrigin
interface PersonRepository extends CrudRepository<Person, Long> {}

```

In the above example CORS support is enabled for the whole `PersonRepository`. `@CrossOrigin` provides attributes to configure CORS support.

```

@CrossOrigin(origins = "http://domain2.com",
    methods = { RequestMethod.GET, RequestMethod.POST, RequestMethod.DELETE },
    maxAge = 3600)
interface PersonRepository extends CrudRepository<Person, Long> {}

```

This example enables CORS support for the whole `PersonRepository` providing one origin, restricted to `GET`, `POST` and `DELETE` methods with a max age of 3600 seconds.

15.7.2. Repository REST Controller method CORS configuration

Spring Data REST fully supports [Spring Web MVC's Controller method configuration](#) on custom REST Controllers sharing repository base paths.

```
@RepositoryRestController
public class PersonController {

    @CrossOrigin(maxAge = 3600)
    @RequestMapping(path = "/people/xml/{id}", method = RequestMethod.GET, produces =
MediaType.APPLICATION_XML_VALUE)
    public Person retrieve(@PathVariable Long id) {
        // ...
    }
}
```

NOTE

Controllers annotated with `@RepositoryRestController` inherit `@CrossOrigin` configuration from their associated repositories.

15.7.3. Global CORS configuration

In addition to fine-grained, annotation-based configuration you'll probably want to define some global CORS configuration as well. This is similar to Spring Web MVC'S CORS configuration but can be declared within Spring Data REST and combined with fine-grained `@CrossOrigin` configuration. By default all origins and `GET`, `HEAD`, and `POST` methods are allowed.

NOTE

Existing Spring Web MVC CORS configuration is not applied to Spring Data REST.

```
@Component
public class SpringDataRestCustomization extends RepositoryRestConfigurerAdapter {

    @Override
    public void configureRepositoryRestConfiguration(RepositoryRestConfiguration config)
    {

        config.getCorsRegistry().addMapping("/person/**")
            .allowedOrigins("http://domain2.com")
            .allowedMethods("PUT", "DELETE")
            .allowedHeaders("header1", "header2", "header3")
            .exposedHeaders("header1", "header2")
            .allowCredentials(false).maxAge(3600);
    }
}
```

Appendix

Appendix A: Using cURL to talk to Spring Data REST

This appendix contains a list of guides that demonstrate interacting with a Spring Data REST service via cURL:

- [Accessing JPA Data with REST](#)
- [Accessing Neo4j Data with REST](#)
- [Accessing MongoDB Data with REST](#)
- [Accessing GemFire Data with REST](#)

Appendix B: Spring Data REST example projects

This appendix contains a list of Spring Data REST sample applications. The exact version of each example isn't guaranteed to match the version of this reference manual.

NOTE To get them all, visit <https://github.com/spring-projects/spring-data-examples> and either clone or download a zipball. This will give you example apps for ALL supported Spring Data projects. Simply navigate to [spring-data-examples/rest](#).

Multi-store example

[This example](#) shows how to mix together several underlying Spring Data projects.

Projections

[This example](#) contains more detailed code you can use to poke around with [projections](#).

Spring Data REST + Spring Security

[This example](#) shows how to secure a [Spring Data REST](#) application in multiple ways with [Spring Security](#).

Starbucks example

[This example](#) exposes 10843 Starbucks coffee shops via a RESTful API that allows to access the stores in a hypermedia based way and exposes a resource to execute geo-location search for coffee shops.